

# INNOVATE

ONLINE CONFERENCE

MACHINE LEARNING  
AND AI EDITION



# Deep dive on Amazon SageMaker

Julien Simon  
Global Evangelist, AI & Machine Learning  
Amazon Web Services

 @julsimon

<https://medium.com/@julsimon>

# Agenda

## 1. Storage

- Amazon S3 & Pipe mode
- Amazon EFS **NEW!**
- Amazon FSx for Lustre **NEW!**

## 2. Training

- Distributed Training
- Managed Spot Training **NEW!**

## 3. Model tuning

## 4. Deployment

- Model compilation: Amazon SageMaker Neo
- Cost optimization: Amazon Elastic Inference



# Storage

# Passing datasets to algorithms

- Amazon SageMaker algorithms accept input data from channels
  - A channel is a **named input source defining a dataset**
  - At least one, up to twenty: Training, validation, test, etc.
- Channel object
  - Name
  - Data source: `S3DataSource` or `FileSystemDataSource`
  - Data format: CSV, RecordIO, etc.
  - Compression type
  - Input mode: `File` or `Pipe` (S3 only)
- The list of channels is passed to CreateTrainingJob
- Amazon SageMaker Python SDK: `Estimator.fit()` receives a dictionary
  - `sagemaker.inputs.s3_input` for S3
  - `sagemaker.inputs.FileSystemInput` for EFS/FSx

# Storing your dataset in Amazon S3

- Simplest option
  - `sagemaker.session.default_bucket()`, `sagemaker.session.upload_data()`
- S3DataSource
  - Location: URI
  - Type: Prefix, manifest, augmented manifest
  - Distribution: Fully replicated (training instances receive the full dataset), or sharded ( $1/n^{\text{th}}$  of the dataset)
- Input mode: **File** or **Pipe**?
  - File: **Copy** the dataset to each training instance (full or  $1/n^{\text{th}}$ )
  - Pipe: **Stream** directly from S3
    - Training **starts faster** and **runs faster**
    - No need to provision lots of **storage** on training instances
    - Train on **arbitrary large datasets**, as they don't need to be fully stored or loaded in RAM any longer
- Pipe mode is supported by **most built-in algorithms** and can be implemented in TensorFlow, **Apache MXNet**, etc.
  - <https://aws.amazon.com/blogs/machine-learning/using-pipe-input-mode-for-amazon-sagemaker-algorithms/>
  - <https://aws.amazon.com/blogs/machine-learning/accelerate-model-training-using-faster-pipe-mode-on-amazon-sagemaker/>





# Demo: A quick look at Pipe mode with TensorFlow

[https://github.com/aws-labs/amazon-sagemaker-examples/tree/master/sagemaker-python-sdk/tensorflow\\_script\\_mode\\_pipe\\_mode](https://github.com/aws-labs/amazon-sagemaker-examples/tree/master/sagemaker-python-sdk/tensorflow_script_mode_pipe_mode)

# Storing your dataset in Amazon EFS

<https://aws.amazon.com/blogs/machine-learning/speed-up-training-on-amazon-sagemaker-using-amazon-efs-or-amazon-fsx-for-lustre-file-systems/>

- If your organization is **sharing data over NFS**, this is a good option
  - Shared datasets, notebooks, etc.
  - Train directly from EFS, **no data movement** required
- Training instances must run in a **VPC**, and open **port 2049** (NFS)
- [FileSystemDataSource](#)
  - Filesystem id: provided by EFS
  - Type: 'EFS'
  - Directory path
  - Access type: Read-only or read-write

```
estimator = TensorFlow(  
    entry_point='tensorflow_mnist/mnist.py',  
    role='SageMakerRole',  
    train_instance_count=1,  
    train_instance_type='ml.c4.xlarge',  
    subnets=['subnet-1', 'subnet-2'],  
    security_group_ids=['sg-1'])  
  
file_system_input = FileSystemInput(  
    file_system_id='fs-1',  
    file_system_type='EFS',  
    directory_path='/tensorflow',  
    file_system_access_mode='ro')  
  
estimator.fit(inputs=file_system_input)
```




# Amazon FSx for Lustre

## Fully managed Lustre file system for compute-intensive workloads


<https://aws.amazon.com/fsx/lustre/>



Massively scalable  
performance



Seamless access to  
your data repositories



Simple  
and fully managed



Native file  
system interface



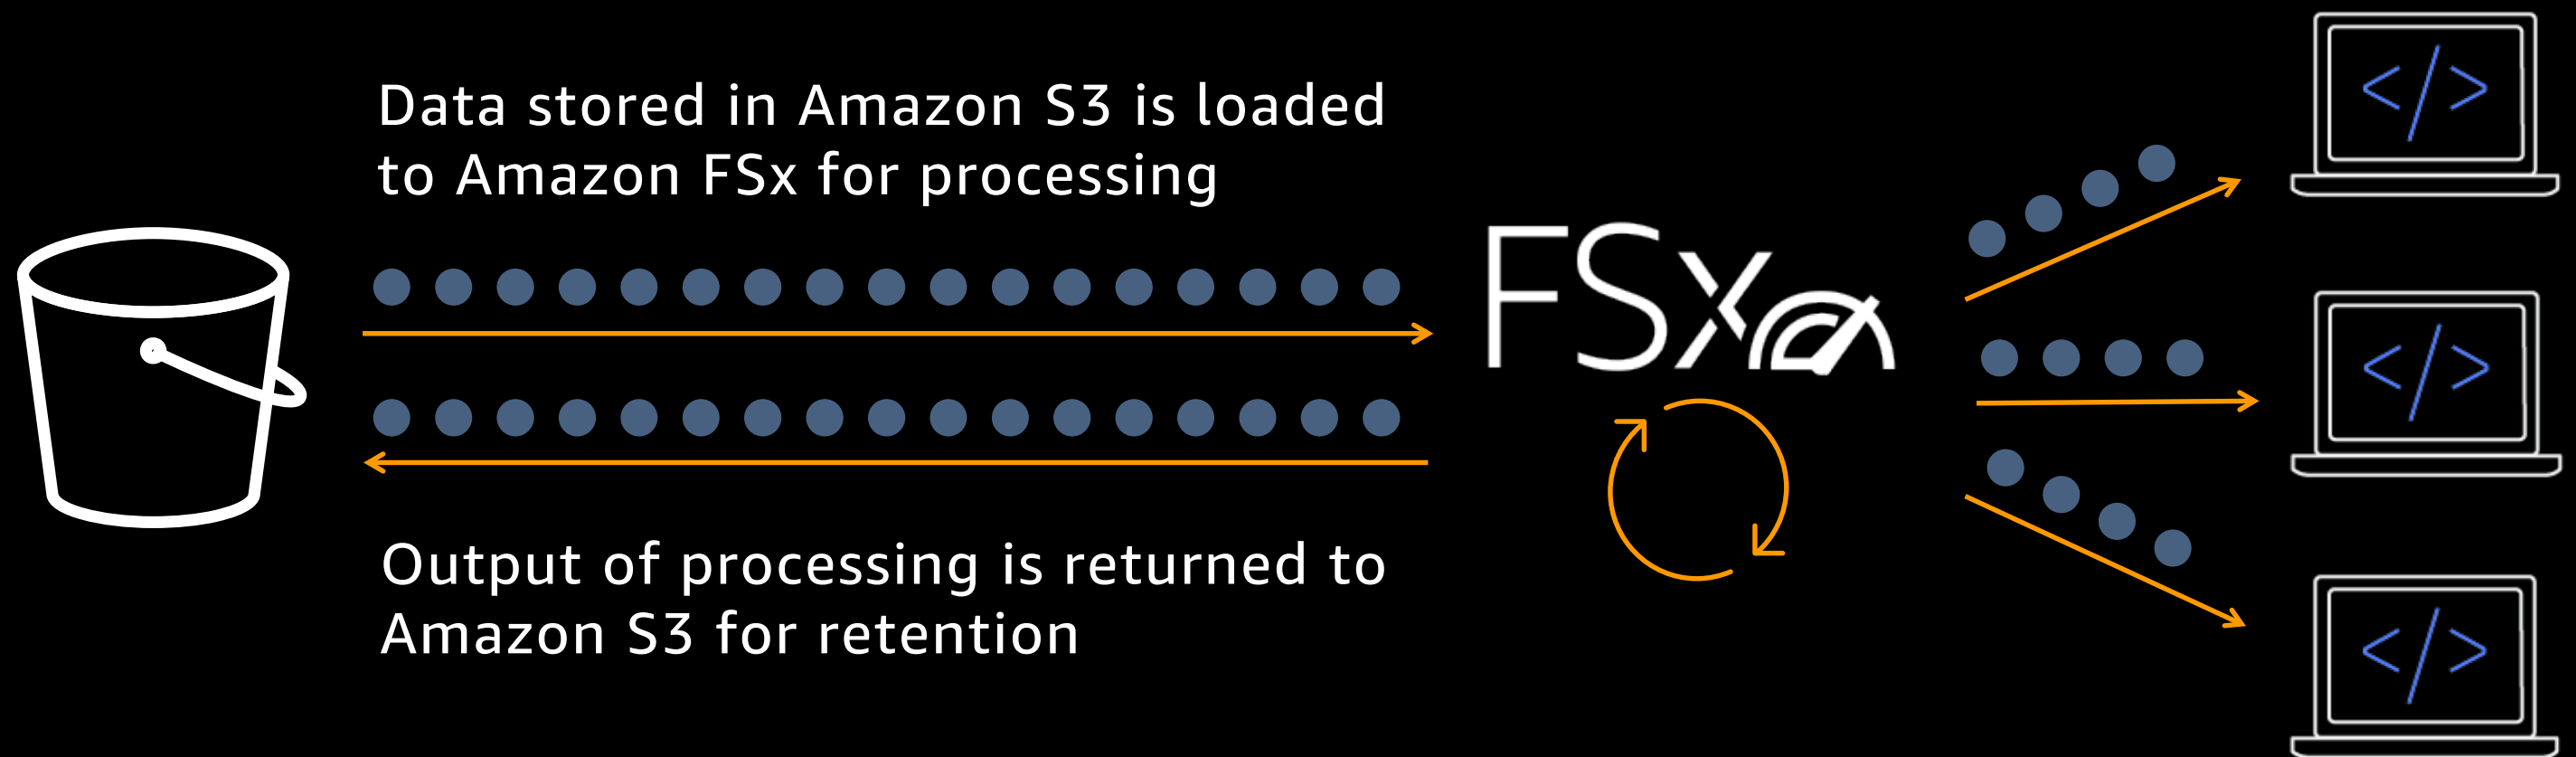
Cost-optimized for  
compute-intensive workloads



Secure  
and compliant

# Seamless integration with Amazon S3

Link your Amazon S3 dataset to your Amazon FSx for Lustre file system, then...



When your workload finishes, simply delete your file system

# Storing your dataset in Amazon FSx for Lustre

<https://aws.amazon.com/blogs/machine-learning/speed-up-training-on-amazon-sagemaker-using-amazon-efs-or-amazon-fsx-for-lustre-file-systems/>

- Best option for **high-performance, low-latency** training
- Create an FSx file system, link it to your S3 bucket, train
- Delete the file system when you're done

- Training instances must run in a **VPC**, and open **port 998** (Lustre)

- **FileSystemDataSource**

- Filesystem id: provided by FSx
- Type: 'FSxLustre'
- Directory path
- Access type: Read-only or read-write

```
estimator = TensorFlow(  
    entry_point='tensorflow_mnist/mnist.py',  
    role='SageMakerRole',  
    train_instance_count=1,  
    train_instance_type='ml.c4.xlarge',  
    subnets=['subnet-1', 'subnet-2'],  
    security_group_ids=['sg-1'])  
  
file_system_input = FileSystemInput(  
    file_system_id='fs-2',  
    file_system_type='FSxLustre',  
    directory_path='/fsx/tensorflow',  
    file_system_access_mode='ro')  
  
estimator.fit(inputs=file_system_input)
```



# Training

# Distributed Training

- Natively available for most built-in algorithms
- Natively available for TensorFlow, Apache MXNet, etc.
- You need to implement it yourself if you use a custom container
- Zoom on TensorFlow: Two modes available
  - **Parameter Server**
    - Asynchronous gradient averaging and weight distribution
    - All instances talk to each other: Networking can become a bottleneck and slow down training
  - **Horovod**
    - Based on Ring-AllReduce algorithm
    - More efficient communication helps scale near-linearly to 256 GPUs
    - <https://github.com/aws-samples/sagemaker-horovod-distributed-training>




# Demo: A quick look at TensorFlow with Horovod

[https://github.com/awslabs/amazon-sagemaker-examples/tree/master/sagemaker-python-sdk/tensorflow\\_script\\_mode\\_horovod](https://github.com/awslabs/amazon-sagemaker-examples/tree/master/sagemaker-python-sdk/tensorflow_script_mode_horovod)




# Amazon EC2 P3dn


<https://aws.amazon.com/blogs/aws/new-ec2-p3dn-gpu-instances-with-100-gbps-networking-local-nvme-storage-for-faster-machine-learning-p3-price-reduction/>



Reduce machine learning  
training time



Better GPU  
utilization



Support larger, more  
complex models

---

## KEY FEATURES

100Gbps of networking  
bandwidth

8 NVIDIA Tesla V100  
GPUs

32GB of memory  
per GPU (2x more)

96 Intel  
Skylake vCPUs  
(50% more than P3)  
with AVX-512

# Managed Spot Training

<https://aws.amazon.com/blogs/aws/managed-spot-training-save-up-to-90-on-your-amazon-sagemaker-training-jobs/>

- Save up to **90%** on training costs
- Fully managed: Obtain spot instances, start training, handle interruptions
- Implement **checkpointing** to resume interrupted jobs
  - Available in built-in algorithms for computer vision
  - Default behavior in TensorFlow
  - If checkpointing is not implemented, the training job is restarted from **scratch**
  - You get billed for data download only once
- CreateTrainingJob
  - `EnableManagedSpotTraining = true`  
`MaxWaitTimeInSeconds`  
`= MaxRuntimeInSeconds`  
`+ time waiting for spot instances`

```
estimator = TensorFlow(  
    entry_point='script.py',  
    role='SageMakerRole',  
    train_instance_count=1,  
    train_instance_type='ml.p3.2xlarge',  
    train_use_spot_instances=true,  
    train_max_run=3600,  
    train_max_wait=7200)  
  
estimator.fit(inputs)
```



# Demo: Fashion-MNIST classification with Keras/TensorFlow

- + Script Mode
- + **Managed Spot Training**
- + Elastic Inference

<https://aws.amazon.com/blogs/machine-learning/train-and-deploy-keras-models-with-tensorflow-and-apache-mxnet-on-amazon-sagemaker/>

<https://gitlab.com/juliensimon/dlnotebooks/tree/master/keras/05-keras-blog-post>

**INNOVATE** | MACHINE LEARNING  
ONLINE CONFERENCE AND AI EDITION

# Tips to speed up training

- Scale out with **distributed training**
- Pick the best format for your dataset
  - Use **protobuf** instead of CSV or JSON
    - <https://github.com/aws/sagemaker-python-sdk/blob/master/src/sagemaker/amazon/common.py>
  - Pack samples into **record-based files**
    - TFRecord (TensorFlow) or RecordIO (MXNet)
    - Splitting in 100MB files looks like the sweet spot
  - Protobuf-encoded + RecordIO ❤️
- Amazon S3: Use **Pipe mode** for large datasets **NEW!**
- Monitor **CPU/GPU usage and network throughput** in Amazon CloudWatch



# Model tuning

# The never-ending quest for hyperparameters

## XGBoost

Tree depth

Max leaf nodes

Gamma

Eta

Lambda

Alpha

...

## Neural networks

Number of layers

Hidden layer width

Learning rate

Embedding dimensions

Dropout

...



# Finding the optimal set of hyperparameters

<https://aws.amazon.com/blogs/machine-learning/amazon-sagemaker-automatic-model-tuning-now-supports-random-search-and-hyperparameter-scaling/>

1. **Manual search:** “I know what I’m doing”
2. **Grid search:** “X marks the spot”
  - Typically training hundreds of models
  - Slow and expensive
3. **Random search:** “Spray and pray”
  - Works better and faster than Grid Search
  - But... but... but... it’s random!
4. **Hyperparameter optimization (HPO):** Use machine learning
  - Requires fewer training jobs
  - Gaussian Process Regression and Bayesian Optimization




# Demo: HPO with Keras

<https://gitlab.com/juliensimon/dlnotebooks/tree/master/keras/04-fashion-mnist-sagemaker-advanced>




# Deployment

**Predictions** drive  
complexity and  
cost in production



# Model optimization is extremely complex



# Amazon SageMaker Neo

<https://aws.amazon.com/blogs/aws/amazon-sagemaker-neo-train-your-machine-learning-models-once-run-them-anywhere/>




Get accuracy  
and performance



Automatic  
optimization



Broad framework  
support



Broad hardware  
support

---

## KEY FEATURES

Integrated with Amazon EC2  
and Amazon SageMaker

Open-source runtime and  
compiler; 1/10<sup>th</sup> the size  
of original frameworks

<https://github.com/neo-ai>



# Compiling ResNet-50 for the Raspberry Pi

## Configure the compilation job

```
{
  "RoleArn": $ROLE_ARN,
  "InputConfig": {
    "S3Uri": "s3://jsimon-neo/model.tar.gz",
    "DataInputConfig": "{\"data\": [1, 3, 224, 224]}",
    "Framework": "MXNET"
  },
  "OutputConfig": {
    "S3OutputLocation": "s3://jsimon-neo/",
    "TargetDevice": "rasp3b"
  },
  "StoppingCondition": {
    "MaxRuntimeInSeconds": 300
  }
}
```

## Compile the model

```
$ aws sagemaker create-compilation-job
--cli-input-json file://config.json
--compilation-job-name resnet50-mxnet-pi

$ aws s3 cp s3://jsimon-neo/model-
rasp3b.tar.gz .

$ gtar tfz model-rasp3b.tar.gz
compiled.params
compiled_model.json
compiled.so
```


## Predict with the compiled model

```
from dlr import DLIRModel
model = DLIRModel('resnet50', input_shape,
                  output_shape, device)
out = model.run(input_data)
```


# Right-sizing your inference infrastructure

- **Statistical ML models, small DL models, dev/test**
  - CPU instances (C5) deliver the best cost/performance ratio
- **Very large DL models**
  - GPU instances (P2 or P3) should work best, especially if you need high throughput
  - If not, C5n could be a reasonable alternative
- But **what about everything in between?**
  - Mid-sized models
  - NLP models
  - Low-throughput, low-latency workloads
  - « Too slow on CPU, not cost-effective on GPU » ?

# Are you making the most of your GPU infrastructure?



Low utilization and high costs




One size does not fit all

# Amazon Elastic Inference


<https://aws.amazon.com/blogs/aws/amazon-elastic-inference-gpu-powered-deep-learning-inference-acceleration/>



Reduce GPU inference  
costs up to 75%



Match capacity  
to demand



Available between 1 and  
32 TFLOPs per accelerator

---

## KEY FEATURES

Integrated with Amazon EC2  
and Amazon SageMaker

Support for TensorFlow  
and Apache MXNet

Single and mixed-precision  
operations



# Demo: Fashion-MNIST classification with Keras/TensorFlow

- + Script Mode
- + Managed Spot Training
- + Elastic Inference

<https://aws.amazon.com/blogs/machine-learning/train-and-deploy-keras-models-with-tensorflow-and-apache-mxnet-on-amazon-sagemaker/>

<https://gitlab.com/juliensimon/dlnotebooks/tree/master/keras/05-keras-blog-post>

**INNOVATE** | MACHINE LEARNING  
ONLINE CONFERENCE AND AI EDITION



# Getting started

<http://aws.amazon.com/free>

<https://ml.aws>

<https://aws.amazon.com/sagemaker>

<https://github.com/aws/sagemaker-python-sdk>

<https://github.com/aws/sagemaker-spark>

<https://github.com/aws-labs/amazon-sagemaker-examples>

<https://gitlab.com/juliensimon/dlnotebooks>

**INNOVATE**  
ONLINE CONFERENCE

MACHINE LEARNING  
AND AI EDITION





# Thank you!

Julien Simon  
Global Evangelist, AI & Machine Learning  
Amazon Web Services

 @julsimon

<https://medium.com/@julsimon>

**INNOVATE** | MACHINE LEARNING  
ONLINE CONFERENCE AND AI EDITION